

Oracle TechTalk: Agent Spec

Run AI Agents Anywhere with Open Agent Specification

Louis Faucon, Database AI Research
13th November 2025

Agenda

- **Background on LLMs and AI Agents**
- Introduction to **Agent Spec**
- **Demo:** Structured Generation
- **Demo:** ReAct agent with Tools
- **Demo:** ReAct agent with MCP

Background

Background – from LLMs to conversational agents

LLMs purely complete missing text; Conversational agents helpfully reply to requests;

Multiple training phases are necessary to produce helpful LLMs for Agents:

- Pretraining: Using very large corpus of text data from the internet
- Prompting: Using System/Agent/User message to produce a conversation.
- Finetuning: Training the model on dedicated data for following instructions
- Alignment: with the preferences of the LLMs developers using Reinforcement Learning from Human Feedbacks




Figure 2 This diagram summarizes our data collection and model training workflow.

Training a Helpful and Harmless Assistant with Reinforcement Learning from Human Feedback, Anthropic <https://arxiv.org/abs/2204.05862>

Background – ReAct Agents

Connecting tools to agents to take actions in the real world

LLMs can be made to decide to invoke external tools:

- System prompt: Contains a description of the tools and how to invoke them
- When receiving a user request, the LLM can then generate either a reply to the user or a Tool Call following a specific format (e.g. JSON)
- After the model generates a Tool Call, the tool is executed and its result appended to the conversation of the Agent as an observation

(1d) ReAct (Reason + Act)

Thought 1: I need to search Cirque du Soleil show Mystere, find the hotel it is in, then find the number of rooms in the hotel.

Act 1: `Search[Cirque du Soleil show Mystere]`

Obs 1: Could not find Cirque du Soleil show Mystere. Similar: ['Cirque du Soleil' ...]

Thought 2: To find the hotel, I can search Mystere (Cirque du Soleil).

Act 2: `Search[Mystere (Cirque du Soleil)]`

Obs 2: Mystère (my-steer) is ... Held in a custom theatre at the Treasure Island Hotel and Casino ...

Thought 3: Mystere is held in a custom theatre at the Treasure Island Hotel and Casino. So I need to search Treasure Island Hotel and Casino next and find the number of rooms in the hotel.

Act 3: `Search[Treasure Island Hotel and Casino]`

Obs 3: Treasure Island Hotel and Casino ... with 2,884 rooms and 220 suites ...

Thought 4: Treasure Island Hotel and Casino has 2,884 rooms and 220 suites. So the answer is 3,104.

Act 4: `Finish[3,104]`


Up-to-date ✓

ReAct: Synergizing Reasoning and Acting in Language Models
<https://arxiv.org/abs/2210.03629>

Background – Model Context Protocol

Standardizing how applications provide tools and context to LLM Agents.

- Open Source standard supported by many systems and platforms
- Simplifies connecting your agents to external services (e.g. web APIs, search engines, ...)
- Provides tools, data sources, and specialized prompts
- MCP servers can either be deployed by application providers or can be run locally by users directly



What is MCP?

<https://modelcontextprotocol.io/docs/getting-started/intro>

Open Agent Specs

Open Agent Specs – Motivation & Context

Proliferation Without Reusability


Proliferation of Agents frameworks:

- 100+ open-source and commercial frameworks
 - No common abstractions for defining agents, tools, and flows
 - Reusing components (e.g., memory modules, planners, toolchains) across frameworks is manual and error-prone

Learnings from Deep Learning & ONNX unification

- Early ML frameworks were siloed: Caffe, Torch, Keras, TensorFlow...
 - Open Neural Network Exchange (ONNX) unified model representation and enabled reusability of models
 - Resulted in a shared ecosystem of reusable models and tooling

→ Open Agent Spec brings a shared foundation that would enable for AI Agents what ONNX did for Deep Learning



Open Agent Specs - A Unified Agent Specification

What is Open Agent Specification (**Agent Spec**)?

“A **framework-agnostic configuration language** for defining AI Agents & Flows.”

Frameworks-agnostic configurations

Agent Spec is a configuration language **independent** from any specific agent framework.

Runtime adapters can be implemented to support Agent Spec **in any agent framework**.

Reproducibility & exchangeability

Every agent defined with Agent Spec can be **executed and evaluated** consistently.

Standardized schemas increase **reproducibility, and exchangeability** across environments.

Modularity & Reusability

Agent Spec provides **pre-built components** and **composition primitives** that make it easy to build, extend, and reuse agents across applications.

Open Agent Specs – Representation and Runtime adapters

Enabling deployment to the right platform

Agent Spec is purely a representation of an agent or agentic workflow

A runtime adapter makes Agent Spec components executable by implementing the behavioral logic defined in the representation

Developers can prototype and develop on one platform (e.g. langgraph) deploy on another

Agent Spec Representation (e.g. RAG Agent)

```
{
  "component_type": "Agent",
  "name": "adaptive_expert_agent",
  "inputs": [{"title": "domain_of_expertise", ...}],
  "outputs": [],
  "llm_config": {
    "component_type": "VllmConfig",
    "name": "Llama 3.1 8B instruct",
    "url": "http://llama.deployment.url/v1",
    "model_id": "Meta-Llama-3.1-8B-Instruct"
  },
  "system_prompt": "You are an expert in {{domain_of_expertise}}. Please help the users with their requests.",
  "tools": [
    {
      "component_type": "ServerTool",
      "name": "rag_tool",
      "description": "Tool that performs RAG",
      "inputs": [{"title": "query", ...}],
      "outputs": [{"title": "results", ...}]
    }
  ],
  "agentspec_version": "25.4.1"
}
```

```
def convert_agspec_llm_to_langgraph(agspec_llm):
  from langchain_openai import ChatOpenAI
...
def convert_agspec_agent_to_langgraph(agspec_agent):
  from langgraph.prebuilt import create_react_agent
...
agent = convert_agspec_agent_to_langgraph(agspec_agent)
```

LangGraph Adapter

```
def convert_agspec_llm_to_autogen(agspec_llm):
  from autogen_ext.models.openai import
  OpenAIChatCompletionClient
...
def convert_agspec_agent_to_autogen(agspec_agent):
  from autogen_agentchat.agents import AssistantAgent
...
agent = convert_agspec_agent_to_autogen(agent_agent)
```

AutoGen Adapter

```
def convert_agspec_llm_to_wayflow(agspec_llm):
  from wayflowcore.models import VllmModel
...
def convert_agspec_agent_to_wayflow(agspec_agent):
  ...
...
agent = convert_agspec_agent_to_wayflow(agspec_agent)
```

WayFlow Adapter

Open Agent Specs – Github & PyPI

Source code and documentation are available on github; Installable directly from PyPI;

```
$ pip install pyagentspec
```

About

Open Agent Spec (Agent Spec) is a framework-agnostic declarative language for defining agentic systems. It defines building blocks for standalone agents and structured agentic workflows as well as common ways of composing them into multi-agent systems.

oracle.github.io/agent-spec/

[Readme](#)
[Apache-2.0, UPL-1.0 licenses found](#)
[Code of conduct](#)
[Contributing](#)

Open Agent Specification
<https://github.com/oracle/agent-spec>

```
$ pip install wayflowcore
```

About

WayFlow is a powerful, intuitive Python library for building sophisticated AI-powered assistants. It is a reference runtime for Agent Spec, with native support for all Agent Spec Agents and Flows.


oracle.github.io/wayflow/

[Readme](#)
[Apache-2.0, UPL-1.0 licenses found](#)
[Code of conduct](#)
[Contributing](#)
[Security policy](#)

WayFlow (runtime for AgentSpec)
<https://github.com/oracle/wayflow>

Open Agent Specs – Code Example – Agents

Code example with pyagentspec



Open Agent Spec – Components – Tools

Supported tool types

Agent Spec makes it easy to integrate into existing software stacks with several types of tools


Tool Type	Description	Agent Spec Tool Configuration
MCP	Interacts with Model Context Protocol (MCP) servers for remote tool execution.	Tool metadata (name, description, inputs, outputs) + URL and communication protocol for the MCP server hosting the tool
Remote	Tool logic runs in an external environment , triggered via RPC or REST.	Tool metadata (name, description, inputs, outputs) + Details of the API request to be performed (e.g. url, payload, ...)
Client	Tool is run by the client , which returns results to the executor.	Tool metadata (name, description, inputs, outputs)
Server	Tool runs in the same runtime as the agent.	Tool metadata (name, description, inputs, outputs) + Requires tool implementation for execution

Open Agent Specs – Components – Flows

Flows specify the steps of a workflow and transitions between them


Control and data flow are constrained by dedicated edge types and define respectively the order of execution of nodes and how data is passed from outputs to inputs of the following nodes

Some nodes might use LLMs, some nodes not (e.g. do REST API call).



Open Agent Specs – Code Example – Flow

Code example with pyagentspec



A Flow is defined with nodes and transitions between the nodes

```

start_node = StartNode(name="start", inputs=[query_property])
tool_node = ToolNode(
    name="hr_lookup_node",
    tool=search_hr_database,
    outputs=[results_property]
)
llm_node = LlmNode(
    name="llm_answer_step",
    llm_config=llm_config,
    prompt_template="Answer the user given the question {{ query }} with the given context: {{ hr_context }}",
    inputs=[query_property, hr_context_property],
    outputs=[answer_property],
)
end_node = EndNode(name="end", outputs=[answer_property])

```

```

control_flow_edges = [
    ControlFlowEdge(name="tool_to_llm",
    from_node=tool_node, to_node=llm_node),
    ...
]

data_edges = [
    DataFlowEdge(name="query_data_edge",
    source_node=start_node, source_output="query",
    destination_node=tool_node, destination_input="query"),
    ...
]

flow = Flow(
    name="HR Flow",
    start_node=start_node,
    nodes=[start_node, tool_node, llm_node, end_node],
    control_flow_connections=control_flow_edges,
    data_flow_connections=data_flow_edges,
)

```

Serialization

Pre-built components have a serialized representation (e.g. JSON)

Flow configurations:

- List of nodes
- Control/data flow definition

Node configurations:

- Inputs/outputs definition
- Per node configuration

LLM configurations:

- Host type, URL
- Model Type
- Generation parameters

Flow information

Control flow definition

Data flow definition

Nodes in flow

```
{
  "component_type": "Flow",
  "inputs": [{"title": "article", "type": "string"}],
  "outputs": [{"title": "summary", "type": "string"}, ...],
  "nodes": [
    {"$component_ref": "start_node_id"},
    {"$component_ref": "llm_node_id"},
    {"$component_ref": "end_node_id"}],
  "control_flow_connections": [
    {
      "component_type": "ControlFlowEdge",
      "from_node": {"$component_ref": "start_node_id"},
      "to_node": {"$component_ref": "llm_node_id"}
    }, ...],
  "data_flow_connections": [
    {
      "component_type": "DataFlowEdge",
      "source_node": {"$component_ref": "start_node_id"},
      "source_output": "article",
      "destination_node": {"$component_ref": "llm_node_id"},
      "destination_input": "article"
    }, ...],
  "$referenced_components": {
    "llm_node_id": {
      "component_type": "LlmNode",
      "inputs": [{"title": "article", "type": "string"}],
      "outputs": [{"title": "summary", "type": "string"}, ...],
      "llm_config": {
        "component_type": "OllamaConfig",
        "default_generation_parameters": {"max_tokens": 512, "temperature": 0.7},
        "url": "http://localhost:11434",
        "model_id": "llama3.1"
      },
      "prompt_template": "Extract the relevant information from the article: {{article}}", ...
    },
    "agentspec_version": "25.4.2"
  }
}
```

Open Agent Specs – Composability and Multi-Agent Systems

From single agents to reusable multi-agent systems

Agent Spec enables developers to **reuse modular components** when building new assistants.

Composability Patterns:

- Flow-in-Flow enables to develop complex logic with well separated reusable subflows
- Agent-in-flow integrates agents into orchestrated workflows
- Flow-in-agent lets agents invoke complex workflow in the same manner tools are invoked
- Swarm / Manager–Worker supports intelligent query routing and specialization

Upcoming support for the **A2A protocol** standardizes communication across agents

Open Agent Specs – Built-in Components

Agent Spec supports a large number of built-in components, and more are being added in future versions

Flow & Node Types

- Flow
- AgentNode
- ApiNode
- BranchingNode
- EndNode
- FlowNode
- InputMessageNode
- LlmNode
- MapNode
- OutputMessageNode
- StartNode
- ToolNode
- ...

Other component Types

- Agent
- OpenAiConfig
- VllmConfig
- OllamaConfig
- OciGenAiConfig
- ClientTool
- ServerTool
- RemoteTool
- MCPTool
- Property
- Swarm
- ...

Open Agent Specs – Evaluation Harness

Evaluation Harness - <https://arxiv.org/pdf/2510.04173v3>

Agent Spec enables comparison of agent and flow patterns across different frameworks.

We ran this evaluation harness on representative tasks from three benchmark datasets:

- LangGraph and WayFlow deliver reliable accuracy and query time stability across these datasets
- AutoGen outperforms on complex reasoning tasks (e.g., on τ_2 - Bench) but with increased latency
- CrewAI's multi-agent coordination yields limited accuracy gains despite longer query times.

Agent Setup	Runtime	F1-score (%)	Time/query (s)
GPT-4.1 mini	-	17.8	1.06 ± 0.41
	WayFlow	79.0	6.22 ± 5.24
ReAct Agent (GPT-4.1 mini)	LangGraph	75.5	4.57 ± 6.21
	AutoGen	67.9	4.64 ± 1.64
	CrewAI	81.1	7.21 ± 6.74
	WayFlow	79.2	4.60 ± 0.92
ReAct Agent (Llama 3.3 70B)	LangGraph	77.6	4.10 ± 0.91
	AutoGen	54.9	3.38 ± 0.31
	CrewAI	82.3	8.12 ± 3.84
Flow-based agentic RAG (GPT-4.1 mini)	WayFlow	88.1	23.19 ± 43.25
	LangGraph	85.3	23.21 ± 55.53

Benchmark results on SimpleQA Verified

Agent Setup	Runtime	EX (%)	Time/query (s)
GPT-4.1 mini	-	50.1	4.26 ± 2.25
GPT-4.1	-	54.4	2.93 ± 4.68
	WayFlow	54.1	7.16 ± 2.72
ReAct Agent (GPT-4.1)	LangGraph	53.7	5.68 ± 3.64
	AutoGen	54.1	6.06 ± 2.97
	CrewAI	54.3	12.20 ± 6.62
NL2SQL Flow (GPT-4.1)	WayFlow	55.5	9.45 ± 5.56
	LangGraph	55.9	7.54 ± 4.11

Benchmark results on BIRD-SQL

Agent Setup	Runtime	Pass ¹ / ^4 (%)	Time/query (s)
	WayFlow	64.7 / 43.0	43.0 ± 13.8
ReAct Agent (GPT-4.1)	LangGraph	71.5 / 51.8	37.5 ± 12.4
	AutoGen	73.5 / 52.6	55.5 ± 17.8
	CrewAI	62.5 / 38.6	62.7 ± 25.4
	WayFlow	60.1 / 32.5	40.9 ± 12.7
ReAct Agent (GPT-4.1 mini)	LangGraph	60.7 / 32.5	38.4 ± 17.4
	AutoGen	58.3 / 36.0	47.4 ± 15.0
	CrewAI	51.5 / 22.8	56.1 ± 25.5

Benchmark results on τ_2 - Bench

Open Agent Specs – Conclusion

Why trying out Agent Spec?

- Framework agnostic
- Sharable agents and flows specification in JSON format
- It's open Source
- Many built-in components that work out of the box
- Support for most types of LLMs models
- Extensive documentations and how-to guides
- Allows you to develop SOTA Agents and Flows

Demos

Demos – Repository

The screenshot shows the GitHub profile of the organization 'LauzHack'. At the top, there's a navigation bar with 'Overview' (selected), 'Repositories 38', 'Projects', 'Packages', and 'People 2'. A search bar is on the right. Below the header, the organization's logo (a red circle with a white cross-like pattern) and name 'LauzHack' are displayed, along with a brief description: 'Student-run hackathon + workshops since 2016 @ EPFL, Switzerland.' It lists 55 followers, location Lausanne, Switzerland, and contact information (website, LinkedIn, Instagram, email). A 'Pinned' section contains two repositories: 'LauzHack.github.io' (HTML, 11 stars, 9 forks) and 'deep-learning-bootcamp' (Jupyter Notebook, 115 stars, 21 forks). Below this is a 'Repositories' section with a search bar and filters for Type, Language, and Sort. Two repositories are listed: 'LauzHack.github.io' (HTML, 11 stars, 9 forks) and 'oracle-agentspec-workshop-2025' (Jupyter Notebook, 1 star, 0 forks, updated 4 hours ago). The repository 'oracle-agentspec-workshop-2025' is highlighted with a yellow box and a yellow arrow points from it towards the text on the right.

You can clone the repository on LauzHack's GitHub to execute the code yourself or just follow along with me.

Conclusion

Conclusion

Want to learn more about AI Agents ?

- Read our [medium blog posts](#)
- Check the docs and guides: <https://oracle.github.io/agent-spec/>

Want to contribute to an Open Source agentdevelopment framework?

- AgentSpec: <https://github.com/oracle/agent-spec>
- Wayflow: <https://github.com/oracle/wayflow>

Want to get more practice and hands-on with AgentSpec & Wayflow?

- If you have registered to the LauzHack Hackathon, come solve Oracle's challenges!

Interested in Internships & Full-time positions at Oracle?

- Send your interests to me: louis.faucon@oracle.com

ORACLE